

# A negotiation meta strategy combining trade-off and concession moves

Raquel Ros · Carles Sierra

Published online: 10 February 2006  
Springer Science + Business Media, Inc. 2006

**Abstract** In this paper we present a meta strategy that combines two negotiation tactics. The first one based on concessions, and the second one, a trade-off tactic. The goal of this work is to demonstrate by experimental analysis that the combination of different negotiation tactics allows agents to improve the negotiation process and as a result, to obtain more satisfactory agreements. The scenario proposed is based on two agents, a buyer and a seller, which negotiate over four issues. The paper presents the results and analysis of the meta strategy's behaviour.

**Keywords** Negotiation

## 1. Introduction

During the last years automated negotiation has become an important challenge in the MAS field. It is the main key for autonomous agent interaction. In a multi agent system we find autonomous agents who decide which actions to execute, when and how. In consequence it is often the case that their own interests conflict with others agents' interests. To solve these conflicts, we must equip them with appropriate negotiation strategies; negotiation is not just a protocol which can be used for agent communication, but also a way to achieve agreements when agents have conflicting interests. To specify a negotiation process we must define [5]:

- Negotiation Protocols: set of rules that govern the interaction (who can participate, which are the negotiation states, what events cause negotiation states to change and what are the valid actions of the participants in each particular state).
- Negotiation Objects: the range of issues over which an agreement must be reached.

---

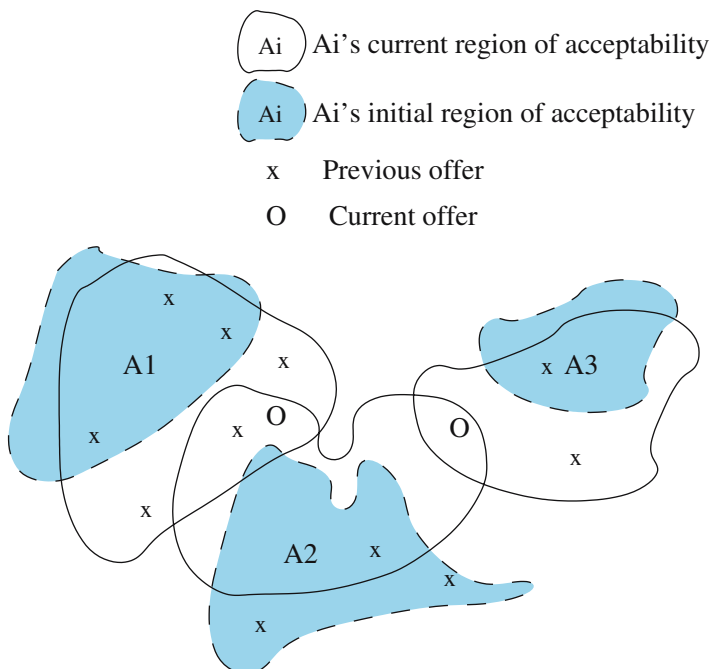
R. Ros · C. Sierra (✉)  
Artificial Intelligence Research Institute—IIIA,  
Spanish Council for Scientific Research—CSIC,  
08193 Bellaterra, Barcelona, Catalonia, Spain  
e-mail: sierra@iiia.csic.es

R. Ros  
e-mail: ros@iiia.csic.es

- Agents' Decision Making Models: the decision making apparatus the participants employ to act in line with the negotiation protocol in order to achieve their objectives. Basically, *how agents negotiate*.

We can define the negotiation process as a *distributed search through a space of potential agreements* (Fig. 1). The dimensionality of the space is determined by the structure of the negotiation object. If we consider each attribute of our negotiation object to have a separate dimension, we clearly see that the space in Fig. 1 concerns two attributes. The preferences of the participants are represented by regions in the negotiation space. If an intersection between these regions exist, then a possible solution to the conflict may be found. The way to reach this solution can be done by interchanging proposals. Formally, a proposal is a *solution to the negotiation problem*. Each proposal can be represented as a point (or region) in the negotiation space. The negotiation process consists of receiving the others agents' proposals and responding to them with a new proposal or an acceptance. The process terminates when the participants find a mutually acceptable point in the negotiation space or when the protocol dictates that the search should be terminated (for whatever reason) without reaching an agreement.

Researchers have proposed different negotiation models. The aim of this work is to combine two existing models in order to improve the negotiation process, i.e. to propose more satisfactory offers for the agents. As a result, we expect to increase the agents' utilities obtained by the agreement achieved. The idea is to switch from one model to the other trying to exploit as much as possible their advantages and to avoid their disadvantages. The first one, we will call it *negoEngine*, is based on concessions [1], and the second one, the *trade-off* strategy [3] where multiple decision variables are traded-off against one another (e.g., paying



**Fig. 1** Negotiation space

a higher price in order to obtain an earlier delivery date or waiting longer in order to obtain a higher quality service). We also propose a modification to the trade-off algorithm in order to improve its performance. Somehow we try to guess the opponent's preferences from the negotiation dialogue in order to propose more acceptable offers. To summarise, this paper presents a step towards the combination of already existing models and a modification of one them to compute more satisfactory offers.

The rest of the paper is organised as follows. Section 2 gives a brief summary of the research done until now on negotiation. Section 3 describes the basic notation used through the paper and the strategies mentioned before. Section 4 details the modification of the trade-off algorithm and introduces the meta strategy proposed. Section 5 presents a scenario for the experiments and the results obtained. An extension for  $n$  negotiating agents is proposed in Section 6. And finally, Section 7 shows the conclusions and future work.

## 2. Related work

During the past years different negotiation approaches have been studied in the fields of game theory and artificial intelligence. Game theory generally assumes agents have complete information about their opponents' preferences. However, in real environments, this situation is unrealistic. For this reason, researches in the AI field have designed new techniques to solve the negotiation problem with incomplete information and uncertainty. Faratin, et al. [1] use heuristic functions to compute the proposals to offer at each time. Parsons, Sierra, and Jennings [12] propose an argumentation model where agents exchange proposals and counter-proposals arguing why they reject an offer. Mugdal and Vassileva [10] propose a sequential decision making in which agents use a preference model of the user incorporating the risk attitude. The decision making is modelled using an influence diagram. Zeng and Sycara [15] propose a sequential decision model which is able to learn. For this purpose, they model the beliefs about the negotiation environment and the participating agents under a probabilistic framework using Bayesian learning representation and updating mechanism. Li, Giampapa, and Sycara [7] study the impact of outside options during a negotiation process. They claim that an outside option affects the negotiation strategy via its impact on the reservation price. Some work has also been done regarding agents with firm deadlines as private information. Fatima, Wooldridge, and Jennings [4] search for the optimal strategy to be selected based on the remaining negotiation time. Sandholm and Vulkan [13] show that the only sequential equilibrium outcome is one where the agents wait until the first deadline, at which point that agent concedes everything to the other. For an extensive review on bilateral negotiation see Li, Giampapa, and Sycara [6]. Research has been mainly focused over different parts of the whole negotiation problem. However, trying to integrate all these parts in one general model is still a task to complete and few proposals can be found in the literature. Among them, Lopes, Mamede, Novais, and Coelho [8] present a generic negotiation model. Their main goal is the integration of two models: an individual behaviour model and a negotiation model based on concessions. As we already mentioned, much work has been done in expanding the negotiation process along different dimensions, for instance, time constraints, outside options, multilateral negotiations, etc. But little work has been done regarding the integration of already designed tactics. In this context, this paper addresses the integration of two negotiation tactics [1,3] in order to improve the outcome. Some modifications have been done on the trade-off algorithm to learn the opponent's preferences in order to compute more satisfactory offers.

### 3. Negotiation strategies

In order to understand the notation used in previous models [1,3] we firstly describe their basics. Then, in the next subsections, we make a quick review of the negotiation models. Let  $i$  ( $i \in \{a, b\}$ ) represent the negotiating agents and  $j$  ( $j \in 1, \dots, n$ ) be the decision variables under negotiation (attributes of our negotiation object). Negotiations can range over quantitative (e.g., price, delivery time, and penalty) or qualitative (e.g., quality of service) decision variables. Quantitative decision variables are defined over a real domain (i.e.,  $x_j^i \in D_j^i = [\min_j^i, \max_j^i]$ ). Qualitative decision variables are defined over a partially ordered set (i.e.,  $x_j^i \in D_j^i = \{q_1, q_2, \dots, q_p\}$ ). Each agent has a scoring function  $V_j^i: D_j^i \rightarrow [0, 1]$  that gives the score it assigns to a value of decision variable  $j$  in the range of its acceptable values. For convenience, scores are kept in the interval  $[0, 1]$ . The relative importance that an agent assigns to each decision variable under negotiation is modelled as a weight,  $w_j^i$ , that gives the importance of decision variable  $j$  for agent  $i$ . We assume the weights of both agents are normalised, i.e.,  $\sum_{1 \leq j \leq n} w_j^i = 1$ , for all  $i \in \{a, b\}$ . An agent's scoring function for a contract,  $\mathbf{x} = (x_1, \dots, x_n)$  in the multi-dimensional space defined by the decision variables' value ranges, is then defined as:  $V^i(\mathbf{x}) = \sum_{1 \leq j \leq n} w_j^i \cdot V_j^i(x_j)$ .

We assume both parties have a deadline by when they must complete the negotiation. This time can be different for each agent and if its deadline passes the agent withdraws from the negotiation. An agent accepts a proposal when the value of the offered contract is higher than the offer the agent is ready to send at that moment in time.

A negotiation thread between agents  $a$  and  $b$  at time  $t_n$  is a finite sequence of proposals from one agent to the other ordered over time:

$$X_{a \leftrightarrow b}^{t_n} = (\mathbf{x}_{a \rightarrow b}^{t_0}, \mathbf{x}_{b \rightarrow a}^{t_1}, \mathbf{x}_{a \rightarrow b}^{t_2}, \dots)$$

Optionally, the last element of the sequence is  $\{\text{accept, reject}\}$ .

#### 3.1. NegoEngine

This subsection describes the first negotiation model (for more details refer to [1]). It is based on defining a set of tactics to be used, either one at a time or as a combination of them. Tactics are the set of functions that determine how to compute the value of a decision variable. For instance:

- *Time dependent*: as time passes, the agent will concede more rapidly trying to achieve an agreement before arriving to the deadline. The value to be uttered by agent  $a$  for a decision variable  $j$  at time  $t$ , with  $0 \leq t \leq t_{\max}^a$  is computed as follows:

$$x_j^a[t] = \begin{cases} \min_j^a + \alpha^a(t)(\max_j^a - \min_j^a) & (1) \\ \min_j^a + (1 - \alpha^a(t))(\max_j^a - \min_j^a) & (2) \end{cases}$$

(1) if  $V_j^a$  is a decreasing function

(2) if  $V_j^a$  is an increasing function

where  $\alpha^a$  is a function depending on time and parametrised by a value  $\beta \in \mathbb{R}^+$ .

$$\alpha^a(t) = \left( \frac{t}{t_{\max}^a} \right)^{\frac{1}{\beta}}$$

For each value of  $\beta$ , infinite number of possible tactics can be represented. However, two qualitatively different classes can be identified: *boulware tactics* if  $\beta < 1$ , and *conceder tactics* if  $\beta > 1$ .

- *Behaviour dependent* or *Imitative*: to imitate the opponent's behaviour.

$$x_j^a[t_{n+1}] = \begin{cases} \min_j^a & \text{if } P \leq \min_j^a \\ \max_j^a & \text{if } P > \max_j^a \\ P & \text{otherwise} \end{cases}$$

The parameter  $P$  determines the type of imitation to be performed. We can find the following families:

- *Relative Tit-For-Tat*: the agent reproduces, in percentage terms, the behaviour that its opponent performed  $\delta \geq 1$  steps ago.

$$P = \frac{x_j^a[t_{n-2\delta}]}{x_j^a[t_{n-2\delta+2}]} x_j^a[t_{n-1}]$$

- *Absolute Tit-For-Tat (Absolute-TFT)*: the same as before, but in absolute terms.

$$P = x_j^a[t_{n-1}] + x_j^a[t_{n-2\delta}] - x_j^a[t_{n-2\delta+2}]$$

- *Averaged Tit-For-Tat (Average-TFT)*: the agent applies the average of percentages of changes in a window of size  $\lambda \geq 1$  of its opponents history.

$$P = \frac{x_j^a[t_{n-2\lambda}]}{x_j^a[t_n]} x_j^a[t_{n-1}]$$

Once we define the tactics to be used during the negotiation process, we also define a combination strategy. We compute the values for the decision variables under negotiation according to each tactic. The final value of each decision variable is a linear combination of these values. To represent this linear combination we use a matrix of weights  $\Gamma$ . Each column represents a tactic and each row, a decision variable. The matrix value  $\gamma_{pm}$  represents the weight assigned to the tactic  $m$  for the decision variable  $p$ . During the negotiation the  $\Gamma$  matrix may change and then, the behaviour of the negotiating agent. The algorithm has linear complexity, as the functions for each tactic are all clearly linear.

*Example* Next, we show an example of a single negotiation step. Suppose agent  $a$  wants to rent an apartment to agent  $b$ . First agent  $a$  must specify its requirements. Then, agent  $b$  initiates the negotiation with its first offer. The decision variables are price (in euros) and area (in  $m^2$ ). The domains and scoring functions are:

$$\begin{aligned} D_{\text{price}}^a &= [250, 700] \\ D_{\text{area}}^a &= [40, 150] \\ V_{\text{price}}^a(x_{\text{price}}) &= \frac{700 - x_{\text{price}}}{450} \\ V_{\text{area}}^a(x_{\text{area}}) &= \frac{x_{\text{area}} - 40}{110} \end{aligned}$$

The tactics being used are the time dependent tactic (TDT) and the behaviour dependent tactic (TFT). For this example we define the  $\Gamma$  matrix as follows:

$$\Gamma_{a \rightarrow b} = \begin{pmatrix} 0.10 & 0.90 \\ 0.15 & 0.85 \end{pmatrix}$$

where the first row corresponds to the price issue, the second one to the area, and the columns to the TDT tactic and the TFT tactic.

Suppose that the negotiation process is in step  $t_2$  and now is agent  $a$ 's turn to make a proposal:

$$X_{a \leftrightarrow b}^{t_2} = \{(700, 40), (272.50, 144.50), (655, 51)\}$$

The new contract to be offered by agent  $a$  at time  $t_3$  is computed in two steps. First we compute the values of the issues according to each tactic. Then, using the  $\Gamma$  matrix, we combine the values suggested by the tactics to obtain the final issue value to offer.

1. TDT:

$$\alpha = \frac{3}{20} = 0.15$$

$$x_{\text{price}}^a[t_3] = 250 + 0.15 \cdot 450 = 317.5$$

$$x_{\text{area}}^a[t_3] = 40 + (1 - 0.15) \cdot 110 = 133.5$$

2. TFT:

$$x_{\text{price}}^a[t_3] = P = \frac{700}{655} \cdot 272.5 = 291.22$$

$$x_{\text{area}}^a[t_3] = P = \frac{40}{51} \cdot 144.5 = 113.33$$

Finally,

$$x_{\text{price}}^a[t_3] = 0.10 \cdot 317.5 + 0.90 \cdot 291.22 = 293.85$$

$$x_{\text{area}}^a[t_3] = 0.15 \cdot 133.5 + 0.85 \cdot 113.33 = 116.35$$

$$\mathbf{x}_{a \rightarrow b}^{t_3} = (293.85, 116.35)$$

### 3.2. Trade-off

The main idea of this tactic is to find a proposal with the same utility as the previous one offered, but expecting to be more acceptable for its opponent (for more details, see [3]). The problem here is how to determinate which offer may increase the opponent's utility, without knowing its preferences. Given an agent  $a$ , who receives a proposal  $\mathbf{y}$  from agent  $b$ , the mechanism should allow agent  $a$  to choose a new proposal  $\mathbf{x}'$  to offer to its opponent which fulfils two conditions:

1. the new proposal  $\mathbf{x}'$  must have the same utility as the offer previously proposed,  $\mathbf{x}$  (this is called  $a$ 's aspiration level);
2. the new proposal  $\mathbf{x}'$  must be the most similar to the offer  $\mathbf{y}$  proposed by  $b$ .

This way, on the one hand we maintain our aspiration level, and on the other hand, we maximise the probability of acceptance of our offer as similarity is in many cases correlated with utility.

Regarding the aspiration level, we define the iso-curves, which are curves formed by all the proposals with the same utility value for an agent:

$$iso_a(\theta) = \{\mathbf{x} | V^a(\mathbf{x}) = \theta\}$$

From this set of proposals, the agent must now choose one. To find the most similar one we use similarity functions which are based on criteria evaluation functions. These evaluation functions determine how much a given element matches the criteria, i.e.,  $h : D \rightarrow [0, 1]$ . Thus a similarity function between two values induced by a single criteria  $h$  can be defined as:  $Sim_h(x, y) = 1 - |h(x) - h(y)|$ . In some cases, multiple criteria can be used to compute

the similarity between two values. To aggregate the individual similarities  $Sim_{h_i}$  a weighted means procedure is employed. Given a domain of values  $D_j$ , a similarity between two values  $x_j, y_j \in D_j$  over  $m$  criteria is defined as:

$$Sim_j(x_j, y_j) = \sum_{1 \leq i \leq m} w_i \cdot (1 - |h_i(x_j) - h_i(y_j)|)$$

where  $\sum_{1 \leq i \leq m} w_i = 1$  is the set of weights representing the importance of the criteria functions in the computation of similarity. For example, suppose we want to determine the similarity between colours [3]:

$$D_{\text{color}} = \{\text{yellow, magenta, violet, green, cyan, red, } \dots\}$$

The similarity can be evaluated based on different criteria: temperature (warm or cold colours), visibility and luminosity. For each criteria, we define an evaluation function (each function represented extensively as sets of pairs):

$$h_t = \left\{ (\text{yellow}, 0.9), (\text{violet}, 0.1), (\text{magenta}, 0.1), \right. \\ \left. (\text{green}, 0.3), (\text{cyan}, 0.2), (\text{red}, 0.7), \dots \right\}$$

$$h_l = \left\{ (\text{yellow}, 0.9), (\text{violet}, 0.3), (\text{magenta}, 0.6), \right. \\ \left. (\text{green}, 0.6), (\text{cyan}, 0.4), (\text{red}, 0.8), \dots \right\}$$

$$h_v = \left\{ (\text{yellow}, 1.0), (\text{violet}, 0.5), (\text{magenta}, 0.4), \right. \\ \left. (\text{green}, 0.1), (\text{cyan}, 1.0), (\text{red}, 0.2), \dots \right\}$$

where for the temperature criteria, warm is 1 and 0 is cold, and for luminosity and visibility, maximum is 1 and minimum is 0. Now, imagine we want to determine the similarity between yellow and red. Given the weights:  $w_t = 0.7$ ,  $w_l = 0.2$ ,  $w_v = 0.1$  for each criteria, the similarity is computed as follows:

$$\begin{aligned} Sim_{\text{color}}(\text{yellow}, \text{red}) &= w_t \cdot (1 - |h_t(\text{yellow}) - h_t(\text{red})|) \\ &\quad + w_l \cdot (1 - |h_l(\text{yellow}) - h_l(\text{red})|) \\ &\quad + w_v \cdot (1 - |h_v(\text{yellow}) - h_v(\text{red})|) \\ &= 0.7 \cdot 0.8 + 0.2 \cdot 0.9 + 0.1 \cdot 0.9 = 0.83 \end{aligned}$$

Thus, we can say that the similarity of colours yellow and red based on these three criteria and their weights is 0.83.


Finally, the similarity between two contracts  $\mathbf{x}$  and  $\mathbf{y}$  over the set of decision variables  $J$  for agent  $a$  is defined as:

$$Sim(\mathbf{x}, \mathbf{y}) = \sum_{j \in J} w_j^a \cdot Sim_j(x_j, y_j)$$

Formalising the trade-off tactic, given the proposal  $\mathbf{x}$  offered by agent  $a$ , and a subsequent offer  $\mathbf{y}$  received from agent  $b$ , where  $\theta = V^a(\mathbf{x})$ , agent  $a$  makes trade-off the following way:

$$trade\text{-}off_a(\mathbf{x}, \mathbf{y}) = \arg \max_{\mathbf{z} \in iso_a(\theta)} \{Sim(\mathbf{z}, \mathbf{y})\}$$

We now explain in more detail the steps performed by the algorithm to compute new proposals. Given the offer  $\mathbf{y}$  proposed by agent  $b$  in time  $t_i$  and the previous proposal  $\mathbf{x}$  offered by agent  $a$  in time  $t_{i-1}$  with  $V^a(\mathbf{y}) < V^a(\mathbf{x})$ , the algorithm must compute a new proposal  $\mathbf{x}'$  to offer in time  $t_{i+1}$ , where  $V^a(\mathbf{x}') = V^a(\mathbf{x})$ . The idea is to increase the utility of the proposal  $\mathbf{y}$ ,  $V^a(\mathbf{y})$ , until it achieves the current aspiration level ( $V^a(\mathbf{x})$ ) after  $S$  steps.



**Fig. 2** Schema of the trade-off algorithm with  $N = 3$  and  $S = 3$

For simplicity, from now on we assume a single step ( $S = 1$ ). As explained in the beginning of the section, the utility of a proposal is the sum of the issues' weighted utilities,  $V_j^a(x_j)$ , under negotiation. Thus, if we increase each individual utility, we also increase the whole proposal's utility. First, the new proposal is initialised,  $\mathbf{x}' = \mathbf{y}$ . Then, the algorithm chooses an issue and increments its utility (either increasing or decreasing its value  $x_j^a$  according to the utility function's monotony). If the aspiration level is not reached yet,  $V^a(\mathbf{x}') < V^a(\mathbf{x})$ , a second issue is chosen and a new value is computed. The process continues until the iso-curve is reached or no issues are left. In the first case, the algorithm finishes and returns the new proposal  $\mathbf{x}'$ . Otherwise, the loop begins again with the first issue.

In summary, the algorithm proposed performs an iterated hill-climbing search in a landscape of possible contracts. The search begins with the last offer received from our opponent and generates a set of  $N$  proposals that lay closer to the iso-curve. At the end of each iteration, the most similar contract is selected. The algorithm terminates when the iso-curve is reached after  $S$  steps. We can see the algorithm steps in Fig. 2. The average time the algorithm takes to complete is linear with respect to the number of decision variables in the negotiation (see [2] for details).

**Example** We now consider the example of a car-dealer to illustrate the model. The decision variables over negotiation are colour, price and delivery time. The first one is a qualitative variable, while the others are quantitative ones. Agent  $b$  (seller) offers its first proposal  $\mathbf{x} = (\text{green}, \text{£}27000, 10 \text{ weeks})$  to agent  $a$  (buyer). Then agent  $a$  proposes its counterproposal  $\mathbf{y} = (\text{yellow}, \text{£}21000, 0 \text{ weeks})$ . To compute the next offer proposed by agent  $b$ , we first need to specify the domains and scoring functions:

$$D_{\text{color}}^b = \{\text{yellow}, \text{violet}, \text{magenta}, \text{green}, \text{blue}, \text{red}\}$$

$$D_{\text{price}}^b = [\text{£}18000, \text{£}35000]$$

$$D_{\text{delivery}}^b = [0 \text{ weeks}, 16 \text{ weeks}]$$



$$\begin{aligned}
V_{\text{color}}^b(x_{\text{color}}) &= \{(yellow, 0.5), (violet, 0.2), (magenta, 0.3), \\
&\quad (green, 0.8), (blue, 0.3), (red, 0.8)\} \\
V_{\text{price}}^b(x_{\text{price}}) &= \frac{x_{\text{price}} - 18000}{17000} \\
V_{\text{delivery}}^b(x_{\text{delivery}}) &= \frac{x_{\text{delivery}}}{16}
\end{aligned}$$

And the weights of the decision variables:  $w_{\text{color}}^b = 0.1$ ,  $w_{\text{price}}^b = 0.8$ ,  $w_{\text{delivery}}^b = 0.1$ . The similarity functions for price and delivery are based on a single criteria: low price (lp) and low delivery (ld):

$$h_{\text{lp}} = \begin{cases} 1 - \frac{x}{40000} & x \in [0, 40000] \\ 0 & \text{otherwise} \end{cases} \quad h_{\text{ld}} = \begin{cases} 1 - \frac{x}{28} & x \in [0, 28] \\ 0 & \text{otherwise} \end{cases}$$

We can now proceed to compute the new offer proposed by agent  $b$ . The seller's aspiration level is  $\theta = V^b(\mathbf{x}) = 0.1 \cdot 0.8 + 0.8 \cdot \frac{27-18}{35-18} + \frac{10}{16} = 0.566$ . While the utility value of agent's  $a$  offer (from the seller's point of view) is  $V^b(\mathbf{y}) = 0.191$ . We have to increase the individual utilities of the decision variables to reach the current iso-curve. We run the algorithm with  $S = 1$  and  $N = 3$ , and we can obtain the following proposals:

$$\mathbf{x}_1 = (yellow, 28301, 5) \quad \mathbf{x}_2 = (red, 28061, 12) \quad \mathbf{x}_3 = (violet, 28673, 7)$$

where all of them have the same utility value as the aspiration level, i.e.,  $V^b(\mathbf{x}_1) = V^b(\mathbf{x}_2) = V^b(\mathbf{x}_3) = \theta$ . The last step is to select the proposal with highest similarity with respect the offer  $\mathbf{y}$  made by agent  $a$ :

$$\begin{aligned}
Sim(\mathbf{y}, \mathbf{x}_1) &= 0.1 \cdot Sim_{\text{color}}(yellow, yellow) + 0.8 \cdot Sim_{\text{price}}(21000, 28301) \\
&\quad + 0.1 \cdot Sim_{\text{delivery}}(0, 5) = 0.1 \cdot 1 + 0.8 \cdot 0.817 + 0.1 \cdot 0.82 \\
&= 0.8356 \\
Sim(\mathbf{y}, \mathbf{x}_2) &= 0.1 \cdot Sim_{\text{color}}(yellow, red) + 0.8 \cdot Sim_{\text{price}}(21000, 28061) \\
&\quad + 0.1 \cdot Sim_{\text{delivery}}(0, 12) = 0.1 \cdot 0.83 + 0.8 \cdot 0.823 + 0.1 \cdot 0.571 \\
&= 0.7985 \\
Sim(\mathbf{y}, \mathbf{x}_3) &= 0.1 \cdot Sim_{\text{color}}(yellow, violet) + 0.8 \cdot Sim_{\text{price}}(21000, 28673) \\
&\quad + 0.1 \cdot Sim_{\text{delivery}}(0, 7) = 0.1 \cdot 0.27 + 0.8 \cdot 0.808 + 0.1 \cdot 0.75 \\
&= 0.748
\end{aligned}$$

Thus, agent  $b$ 's next offer is  $\mathbf{x}_1$ . That is,  $\mathbf{x}' = (yellow, £28301, 5 \text{ weeks})$ .

## 4. Contributions

Next we explain the contributions of this work. First we detail the modification the trade-off algorithm and then we explain the meta strategy proposed.

### 4.1. Modification of the trade-off algorithm

After understanding the behaviour of the trade-off algorithm, we can notice that the order in which each issue is selected determines the final outcome. Given the list of decision variables, the first ones have a higher probability to be modified than the last ones. Thus, in order to propose more satisfactory offers to our opponent, we propose to order the issues according to the opponent's guessed preferences. During human negotiation, it is easy to see that the most important variables are the ones with less variations between offers. For example, if we are not interested in delivery time, it makes no difference to us to change it as our opponent demands it. But, in the case of the price issue, it is important to us to try to keep it as stable as possible with small variations. Then, from our opponent's contracts history we can deduce

somehow its own preferences. With this information we can propose more satisfactory offers varying first the values of those issues that are not so important to our opponent. If we order the decision variables following our opponent's preferences (first those less preferred), the algorithm will begin modifying those that are not so important. In the best case, if the new proposal already achieves the current aspiration level, no more changes will be needed and the most preferred issues may maintain their original values. In the worst case, all issues will be modified. But even in this case, the utility gain needed to achieve the current level will be lower when the most preferred issues are computed.

As a summary, we use the similarity approach presented in Faratin et al. [3] but using as much as possible the knowledge about our opponents' preferences. We bias the exploration in the similarity landscape.

We then define the *variability* of a decision variable in a window of size  $m$  of the contracts history as:

$$f(j) = \frac{\sum_{i=0}^{m-2} |x_j^i[t_{n-2i}] - x_j^i[t_{n-2(i+1)}]|}{(m-1) \cdot \Delta x_{\max}^j}$$

with  $\Delta x_{\max}^j = \max(D_j^i) - \min(D_j^i)$ ,  $m > 1$  and  $t$  the current time. The resulting algorithm includes the computation of the issues' variability and the storing of the offers proposed by the opponent. Thus, the main steps are:

#### Algorithm Smart Trade-off

1. Store received proposal  $\mathbf{y}$  in the contract history
2. **For** each decision variable  $i$  **do**  
    Compute\_variability( $i$ )
3. Order the decision variables based on their variability
4. Compute a new offer using the trade-off algorithm

#### 4.2. Meta strategy

First, we review the advantages and disadvantages of both models. On one hand, the *negoEngine* tactic allows us to compute offers considering the remaining time to end the negotiation process and our opponent's behaviour, both important aspects to consider when trying to achieve an agreement. The disadvantage of this model is that every offer proposed is a concession; this means that our aspiration level decreases in every step of the negotiation process. On the other hand, the trade-off algorithm advantage is that it searches all possible offers which maintain our aspiration level. Thus, our utility gain does not decrease during the negotiation until it is deliberately indicated. An external mechanism is defined to decrease the current aspiration level to achieve an agreement (otherwise, if we never concede, the chance of achieving an agreement is minimum). Faratin et al. proposed to decrease the aspiration level by a predefined amount whenever a deadlock was detected. The problem is that other aspects, as time, are not taken into account.

After reviewing the advantages and drawbacks of the models, we now proceed to describe the meta strategy designed. The main idea is to exploit as much as possible the current aspiration level. If no agreement is reached in a given negotiation step, we reduce our aspiration level expecting to find, in a lower level, a new proposal that satisfies both participants. To manage this behaviour the agent applies a trade-off tactic to maintain the aspiration level until a deadlock is achieved. A deadlock is detected when the last offer proposed by the opponent

does not improve the utility of the offer proposed two steps before. Then, the negoEngine tactic is used in order to decrease the current aspiration level. Using this strategy ensures us to concede in a more rational way, considering the remaining time to end the negotiation and our opponent's behaviour. Next example shows the meta strategy behaviour from the initial state until a deadlock situation is detected:

$t$	$V^a(\mathbf{x})$	$V^a(\mathbf{y})$
$t_0$	0.800	
$t_1$		0.200
$t_2$	0.800	
$t_3$		0.334
$t_4$	0.800	
$t_5$		0.329
$t_6$	0.751	

where  $V^a(\cdot)$  is agent  $a$ 's utility function,  $\mathbf{x}$ , agent  $a$ 's proposals, and  $\mathbf{y}$  corresponds to agent  $b$ 's proposals.

The initial aspiration level is set to 0.8. At time  $t_0$  agent  $a$  proposes an offer. Then, agent  $b$  offers a proposal in time  $t_1$  with a utility value of 0.2 for agent  $a$ . The process continues until time  $t_5$ , where  $b$ 's utility proposal decreases compared to the proposal received in time  $t_3$ . The meta strategy detects the deadlock situation. Thus, in time  $t_6$  agent  $a$  computes the new proposal using the negoEngine tactic, decreasing the current aspiration level to 0.751.

#### Algorithm Meta Strategy

1. **While** deadline is not reached,  $t_{\max}$ , **or** no agreement is found,  $V^a(\mathbf{y}) < V^a(\mathbf{x})$ ,  
**do**
  - (a) Given the last offer  $\mathbf{x}$  proposed by agent  $a$ , compute  $\theta$   
 $\theta = V^a(\mathbf{x})$
  - (b) **If** no deadlock **then** propose a new offer  $\mathbf{x}'$  using the smart trade-off tactic.  
**else** propose a new offer  $\mathbf{x}'$  using the negoEngine tactic.
2. **If** the deadline  $t_{\max}$  is reached **then** withdraw and terminate.  
**Else** accept the proposal  $\mathbf{y}$  and terminate.

As we already mentioned, the complexity of each tactic is linear with respect to the number of decision variables. Thus, combining both tactics does not increase the overall complexity of the meta strategy presented. Regarding the study of the variability of the issues, for each issue, we have to compare  $m - 1$  values from the window of size  $m$ . Thus, the number of operations is  $n(m - 1)$ .

## 5. Experiments

In this section we first present the scenario used in our experimentation. Then the experiments realized are explained, and finally we proceed to the analysis of the results obtained.

The experiments involve two players,  $a$  and  $b$  bargaining over fabric products. The decision variables under negotiation are colour, material, price and delivery time. Even though colour and material issues are discrete decision variables, to simplify the scenario description we assume that all are modelled as a continuous domain. Regarding the colour issue, values are ordered based on colour *temperature* (meaning 0 for extreme cold colours and increasing with the warmth of the colour). The same way, material issue is based on *wrinkle resistance*

(0 means no wrinkle resistance at all, increasing as more resistant is the material).

$$D_c = [0, 5]$$

$$D_m = [0, 4]$$

$$D_p = [30 \text{ euros}, 70 \text{ euros}]$$

$$D_d = [5 \text{ days}, 15 \text{ days}]$$

The weight vectors representing the agents' preferences for each decision variable are fixed during the negotiation:

$$W^a = [0.35, 0.15, 0.45, 0.05] \quad W^b = [0.10, 0.15, 0.40, 0.35]$$

where each weight corresponds to colour, material, price and delivery time issues. Regarding to the evaluation functions we use linear functions:

$$\begin{aligned} V_c^a(x) &= \frac{x}{5} & V_c^b(x) &= \frac{5-x}{5} \\ V_m^a(x) &= \frac{x}{4} & V_m^b(x) &= \frac{4-x}{4} \\ V_p^a(x) &= \frac{70-x}{70-30} & V_p^b(x) &= \frac{x-30}{70-30} \\ V_d^a(x) &= \frac{15-x}{15-5} & V_d^b(x) &= \frac{x-5}{15-5} \end{aligned}$$

And finally, the similarity functions shared by both agents. Similarity for price and delivery are each based on two criteria: low and high price,  $h_{lp}$  and  $h_{hp}$  respectively; and fast and slow delivery time,  $h_{fd}$  and  $h_{sd}$ . The weights are defined as follows:  $w_{lp}^a = 0.8$ ,  $w_{hp}^a = 0.2$ ,  $w_{fd}^a = 0.8$  and  $w_{sd}^a = 0.2$ , for agent  $a$ ; and  $w_{lp}^b = 0.2$ ,  $w_{hp}^b = 0.8$ ,  $w_{fd}^b = 0.2$  and  $w_{lp}^b = 0.8$ , for agent  $b$ .

$$\begin{aligned} h_{hp}(x) &= \begin{cases} 1 & x > 100 \\ \frac{x-20}{80} & x \in [20, 100] \\ 0 & x < 20 \end{cases} \\ h_{fd}(x) &= \begin{cases} 1 & x < 5 \\ \frac{30-x}{25} & x \in [5, 30] \\ 0 & x > 30 \end{cases} \\ h_{lp}(x) &= \begin{cases} 1 & x < 20 \\ \frac{100-x}{80} & x \in [20, 100] \\ 0 & x > 100 \end{cases} \\ h_{sd}(x) &= \begin{cases} 1 & x > 30 \\ \frac{x-5}{25} & x \in [5, 30] \\ 0 & x < 5 \end{cases} \end{aligned}$$

Colour and material similarity are represented as linear function based on a single criteria:

$$h(x) = \frac{x - \min}{\max - \min}$$

The experiments involve a complete negotiation process. An agent makes its first offer and the opponent responds with another one. The interaction continues until an agreement is found or the negotiation time expires. The first offer proposed by the agents is computed with the negoEngine tactic and then different combinations of tactics are used to compare the performance of our meta strategy. Thus, we define the next types of agents:

- *NegoTO agent*: this agent employs the meta strategy defined on Section 4.2. That is, applying the trade-off tactic until it reaches a deadlock, and then making an offer with the negoEngine tactic.

- *Random agent*: the next strategy to compute the new offer is chosen randomly. For instance: negoEngine, trade-off, trade-off, negoEngine, trade-off, negoEngine, negoEngine, ...
- *Alternate agent*: altering both strategies throughout the negotiation process, one at a time: negoEngine, trade-off, negoEngine, trade-off, negoEngine, trade-off, ...
- *TO agent*: in this case, the agent only applies the trade-off tactic while the utility of the offer received is higher than the previous one received. Otherwise, the aspiration level is decreased by a fixed 0.05 value and a new proposal is generated.
- *Nego agent*: this agent only uses the negoEngine tactic during the negotiation.

Regarding the negoEngine tactic, we model five types of behaviours: *very bouldware* ( $B$ ), *bouldware* ( $b$ ), *neutral* ( $n$ ), *conceder* ( $c$ ) and *very conceder* ( $C$ ). We also combine two tactics: time dependent and relative tit-for-tat. The weights of the  $\Gamma$  matrix are  $\gamma_{td} = 0.1$  and  $\gamma_{tft} = 0.9$  respectively.

Two measures were obtained during the experimentation:

1. *utility product*: once an agreement is achieved, the product of the utilities obtained by both participants is computed.
2. *utility difference*: once an agreement is achieved, the of the utilities obtained is computed.

The first measure indicates us the joint outcome, while the second one, indicates the distance between both utilities. There is an important relation between these two measures and compromise should be taken in account. Even though a high joint outcome is expected, it is also important that the difference between both utilities is low. For example, an agent can obtain an utility of 0.75, while the other one, 0.48. The joint utility would be 0.36, which is quite good. But the difference is also high enough to consider the contract as an unsatisfactory agreement (with an assumption of equal negotiation power). For this reason, we will evaluate the results obtained, not only based on the utility product, but also on the utility difference.

We realized 100 bilateral negotiations for every pair of agents to obtain an average of the outcome utility (each execution will variate as the trade-off algorithm includes randomness): *NegoTO agent* versus *NegoTO agent*, *NegoTO agent* versus *Random agent*, *Alternate agent* versus *NegoTO agent*, *Alternate agent* versus *Random*, and so on. We also modified the behaviour of every pair, changing from a *very conceder* one, to a *very bouldware* to study its influence on the final outcome. The negotiation deadline was fixed to 40 steps for both agents. The next tables show the averages of the utility outcomes for both agents ( $V^a(\mathbf{x})$  corresponds to the utility obtained by the reference agent indicated in the table's caption and  $V^i(\mathbf{x})$ , to the utility obtained by the rest of the agents  $a_i$ ), the utility products and the utility differences obtained with *neutral behaviour* in all cases.

We can clearly see that in general the *NegoTO agents* improve the negotiations achieving satisfactory agreements for both participants. As expected, negotiations performed by at least one *NegoTO agent* fulfil the desired properties, i.e., high utility products and low utility differences. As we can see on Table 1(a) the highest utility product (0.360) is obtained between a *NegoTO agent* and a *TO agent*, but the utility difference (0.244) also increases. This means that while the *NegoTO agent* achieves a higher utility, its opponent achieves a lower one. The best equilibrium is found when two *NegoTO agents* negotiate together (0.350 for the utility product and 0.039 for the utility difference). Also notice that comparing to the rest of the agents' utilities, the *NegoTO agent* achieves in all cases, the higher one ( $V^a(\mathbf{x}) > V^i(\mathbf{x})$ ), where  $a$  is the *NegoTO agent*, and  $i$ , all the rest).

In Table 1(b), where the reference agent is the *Alternate agent* again we confirm the results shown before. The *NegoTO agent* achieves the higher utility products and the lower utility

**Table 1** Where \* refers to the utility product, and  $|-|$ , to the utility difference. (a) utility measures obtained by a *NegoTO agent* versus all the rest. (b) Utility measures obtained by a *Alternate agent* versus all the rest. *Neutral* behaviour in all cases

$agent_i$	$V^a(\mathbf{x})$	$V^i(\mathbf{x})$	*	$ - $
(a)				
NegoTO	0.611	0.572	0.350	0.039
Random	0.649	0.514	0.333	0.135
Alternate	0.634	0.514	0.326	0.120
TO	0.734	0.490	0.360	0.244
Nego	0.742	0.303	0.224	0.439
(b)				
NegoTO	0.562	0.592	0.332	0.030
Random	0.592	0.553	0.327	0.039
Alternate	0.608	0.543	0.330	0.065
TO	0.658	0.512	0.337	0.146
Nego	0.630	0.399	0.252	0.231

differences. While the *Alternate agent*, except for the *NegoTO agent*, always obtains a higher utility compared to its opponents' utilities. This means that competing with other agents, the alternate meta strategy does a quite good performance obtaining advantages among the rest.

Regarding the *Random agent* the best combination is obtained when negotiating against a *TO agent*. But as depicted in Table 2(a) comparing the utility obtained in the agreement against the *NegoTO agent* and the *Alternate agent*, it finalises the negotiation with a lower utility gain. Thus, we can say that it cannot improve the other agents performance. In Table 2 (b) we observe that the *TO agent* achieves the higher utility product against itself, and also the lower utility difference, but it cannot improve the others final utility gain, except when negotiating with a *Nego agent*. Finally, and also as expected, the *Nego agent* is the one with the worst performance. This is obvious as it has a concession strategy where it does not look for an improvement. It only tries to achieve an agreement as soon as possible and conceding as much as it can. We can see the results in Table 2(c).


Figure 3 depicts two examples of a complete negotiation process. We represent the proposals offered by the buyer (*NegoTO agent*) in circles and the ones offered by the seller ( $agent_i$ ) in crosses. The  $x$  axis represents the utility perceived by the seller, while the  $y$  axis, represents the buyer's utility (in the range  $[0,1]$ ). After analysing the complete process between a *NegoTO agent* and the rest of the agents, we can observe two situations:

- in the best case, the *NegoTO agent* tries to maintain the current utility gain as much as it can. If its opponent always concedes, offering better proposals at each step, the *NegoTO agent* continues sending proposals computed with the trade-off algorithm as long as the negotiation process lasts. Finally the agreement ends without modifying (or modifying very little) the initial aspiration level.
- in the worst case, the *NegoTO agent* behaves similar to an alternate strategy. Suppose we have a *very bouldware* opponent which makes few concessions. In the beginning the *NegoTO agent* proposes an offer using the trade-off algorithm. Next a deadlock is detected (because the offers received do not improve previous ones), and a new proposal is generated with the negoEngine tactic. In the next step it tries again with the trade-off tactic. Then, if a new deadlock occurs, the negoEngine tactic is employed one more time. The process is repeated again and again until an agreement is reached or agents withdraw.

**Table 2** Where \* refers to the utility product, and  $|-|$ , to the utility difference. (a) utility measures obtained by a *Random agent* versus all the rest. (b) utility measures obtained by a *TO agent* versus all the rest. (c) Utility measures obtained by a *Nego agent* versus all the rest. *Neutral* behaviour in all cases

$agent_i$	$V^a(\mathbf{x})$	$V^i(\mathbf{x})$	*	$ - $
(a)				
NegoTO	0.543	0.613	0.333	0.070
Random	0.576	0.558	0.321	0.018
Alternate	0.550	0.574	0.316	0.024
TO	0.598	0.562	0.336	0.036
Nego	0.637	0.407	0.259	0.230
(b)				
NegoTO	0.437	0.776	0.339	0.339
Random	0.562	0.606	0.340	0.044
Alternate	0.503	0.638	0.321	0.135
TO	0.636	0.565	0.360	0.071
Nego	0.579	0.453	0.262	0.127
(c)				
NegoTO	0.341	0.728	0.248	0.387
Random	0.483	0.591	0.286	0.109
Alternate	0.423	0.605	0.256	0.183
TO	0.484	0.600	0.290	0.115
Nego	0.506	0.494	0.250	0.011

It is also important to mention that changing the behaviour of the negoEngine tactic did not really affect the final outcome. The reason is that the participation of this strategy during the negotiation process is mainly used to finish the execution on time, and not to improve the final outcome. As more *conceder* is the behaviour, the faster the agreement is achieved. In situations where the negotiation time is significant to evaluate the performance of the negotiation, the agent's behaviour must be taken into account. If a time cost is introduced into the model, the negoEngine parameters would need to be modified in order to reduce the necessary time to reach an agreement, allowing maybe greater concessions in the proposals computed. As the trade-off algorithm does not take into account the time factor the meta strategy should also be modified to switch from one strategy to the other one when time cost increases.

**Fig. 3** Negotiation process between two players. On the left, *NegoTO* versus *NegoTO*; and on the right, *NegoTO* versus *Alternate*

More experiments were realized modifying other parameters as the issues' weights,  $w_j^i$ , the domains,  $D_j^i$ , and utilities functions,  $V_j^i$ . They did not caused a significance variation on the results shown.

## 6. Extending the negotiation to $n$ agents

Next we propose an extension of the negotiation process in order to consider  $n$  negotiating agents instead of only two. The offers are generated using the tactics explained in this paper either with the concession tactics or the trade-off algorithm. For this purpose, three aspects have to be considered. First, how  $n - 1$  proposals from  $n - 1$  agents may affect the behaviour of the concession tactics. Second, how to extend the trade-off algorithm to handle  $n - 1$  proposals. And third, we must define a new negotiation protocol for  $n$  agents.


### 6.1. Concessions tactics

As showed in Section 3.1 different tactics can be used to generate a new offer. Some tactics are not influenced by the proposals received during the negotiation process. Thus, they are not affected by the fact that the negotiation is being performed between two agents or more. An example are the *time-dependent tactic* or the *resource-dependent tactic*. They generate new proposals with independent information, not taking into account the offers received. But there is also another family of tactics that depend on the proposals sent by the other agent (since we have been talking about a bilateral negotiation until now). We refer to the *behaviour-dependent tactics*. These tactics generate new proposals depending on how the other agent behaves, i.e., “if he concedes, then I concede.” In some way, they try to model the opponent's behaviour. We believe that in a situation where  $n$  agents negotiate, these tactics could model a social behaviour. Even more, they could detect coalitions between agents as they all will perform similar behaviours. Instead of considering the proposals offered by a single agent in  $\delta$  steps ago, each tactic may compute the new proposal based on the offers proposed by all agents.

### 6.2. Trade-off algorithm

In an scenario with two agents,  $a$  and  $b$ , given an offer  $y$  sent by agent  $b$ , agent  $a$  computes a counterproposal,  $x'$ , with the same utility as the previous one offered,  $x$ , and as similar as possible to  $y$  using the trade-off algorithm. The new scenario proposed is as follows. Given  $n - 1$  agents, each of them offers a proposal,  $y_j$ , to agent  $a$ . Thus, the algorithm has to compute a new counterproposal fulfilling the same two conditions: it must have the same utility value as the previous one, and it has to be the *most* similar to the  $n - 1$  proposals received. Figure 4 illustrates the new situation. The white circles represent the  $n - 1$  offers from the agents; the black circle corresponds to agent  $a$ 's previous offer; the dashed circles represent the  $P$  possible counteroffers with same utility value  $\theta$  (from now on, we talk about  $P$  children instead of  $N$  just to avoid confusions with the  $n$  agents involved in this new approach); and finally, the gray circle is the selected offer to be sent to the agents. Regarding the first condition, the behaviour of the algorithm remains the same. It computes  $P$  proposals, each with the same utility value  $\theta$ . At this point, one has to be chosen to be offered as a counterproposal





**Fig. 4** Schema of the trade-off algorithm between  $n$  agents

(second condition). For this purpose, we propose to select  $\mathbf{x}'$  as the one with minimum square difference respect the  $n - 1$  offers:

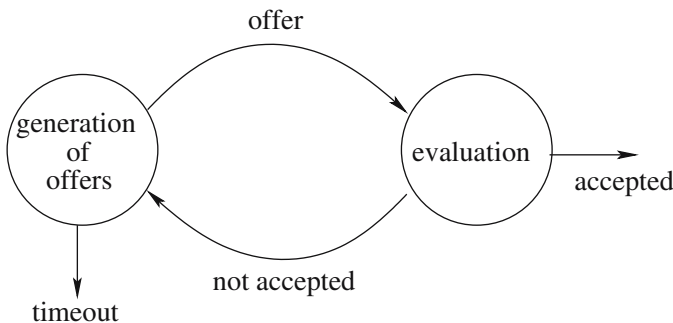
$$\mathbf{x}' = \arg_{\mathbf{x}_i \in \{\mathbf{x}_1, \dots, \mathbf{x}_p\}} \min \left\{ \sum_{j=1}^{n-1} (1 - \text{Sim}(\mathbf{x}_i, \mathbf{y}_j))^2 \right\}$$

### 6.3. Protocol

The negotiation protocol for  $n$  agents is as follows. An agent proposes a first offer to the rest of the  $n - 1$  agents. Next, all respond either with an *accept* or a *reject*. If all accept, then the process finishes with the last offer proposed. Otherwise, another agent (or the same one) generates a new offer. If no agreement is reached until a fixed deadline,  $t_{\max}$ , the process finishes. Figure 5 shows possible states and transitions of the negotiation process presented. Also a token-based protocol must be defined to decide who generates the next offer.

## 7. Conclusions and future work

This paper presents the design of meta strategies to combine negotiation tactics. More precisely, a model based on concession functions and a trade-off tactic. As we have seen, the



**Fig. 5** Negotiation protocol for  $n$  agents

first ones try to achieve an agreement decreasing the expected utility of the final proposal accepted by the participants, while the latter searches a satisfactory proposal maintaining the utility as much as possible. Combining tactics allows agents to better adapt to different situations. In this article we described two simple combinations (random and alternate), and a more accurate one, the meta strategy showed on Section 4. After the experiments we could see that the meta strategy developed obtains better results than the other combinations. In the worst case it behaves as an alternate strategy, which also performs quite well. In the best case, it exploits as much as possible the trade-off tactic in order to maintain the current aspiration level. We also presented a mechanism to detect our opponent's preference in order to propose more satisfactory offers. As a consequence, the probability of acceptance increases.

As future work we propose the refinement of the parameters involved in both models. For this purpose we suggest genetic algorithms due to the huge quantity of parameters of the models. Some of them have been already tested, as for example, the behaviour of the agents when using the concession tactics. But still other parameters should be tested in order to investigate their optimal values and to determine how they affect the negotiation process. It would be also interesting to include other negotiation models, such as argumentation based models [12]. In these models agents respond not only with a new proposal, but also with an argument explaining why they produced that offer or why they reject the received one. This way, agents can get to know its opponents' interests and thus propose them more satisfactory offers. More modifications can be experimented, as in Teuteberg [14], where fuzzy-logic is used to model the utility function. Also, bilateral negotiation could be transformed into a multilateral negotiation. Instead of performing negotiations between two agents, an agent could negotiate against more than one agent. We presented the first steps towards this approach in Section 6. Finally, including time restrictions to achieve an agreement could be done. In this paper, time influence is only represented during the execution of the negoEngine tactic, where the new proposal is computed based on a *time dependent* tactic. It would be interesting to introduce a cost function in the meta strategy to consider time as the negotiation process progresses.

## References

1. Faratin, P., Sierra, C., & Jennings, N. R. (1998). Negotiation decision functions for autonomous agents. *Robotics and Autonomous Systems*, 24, 159–182.
2. Faratin, P., Sierra, C., & Jennings, N. R. (2000). Using similarity criteria to make negotiation trade-offs. *Proceedings of the 4th. International Conference on Multi-Agent Systems* (pp. 119–126). Boston, USA.
3. Faratin, P. Sierra, C., & Jennings, N. R. (2002). Using similarity criteria to make issue trade-offs in automated negotiations. *Artificial Intelligence* 142, 205–237.
4. Fatima, S., Wooldridge, M., & Jennings, N. R. (2001). Optimal negotiation strategies for agents with incomplete information. In *Proceedings 8th International Workshop on Agent Theories, Architectures and Languages (ATAL)* (pp. 53–68). Seattle, USA.
5. Jennings, N. R., Faratin, P., Lomuscio, A. R., Parsons, S., Sierra, C., & Wooldridge, M. (2001). Automated negotiation: prospects, methods and challenges. *International Journal of Group Decision and Negotiation*, 10(2), 199–215.
6. Li, C., Giampapa, J., & Sycara, K. (2003). A review of research literature on bilateral negotiations. In Tech. report CMU-RI-TR-03-41, Robotics Institute, Carnegie Mellon University, November.
7. Li, C., Giampapa, J., & Sycara, K. (2004). Bilateral negotiation decisions with uncertain dynamic outside options. In *The First IEEE International Workshop on Electronic Contracting*, IEEE Computer Society (pp. 54–61). Los Alamitos, California (USA).
8. Lopes, F., Mamede, N., Novais, A.Q., & Coelho, H. (2002). A negotiation model for autonomous computational agents: Formal description and empirical evaluation. *Journal of Intelligent and Fuzzy Systems*, 12, 195–212.

9. Matos, N., & Sierra, C. (1999). Evolutionary computing and negotiation agents. *Agent Mediated Electronic Commerce*, 1571, 126–150.
10. Mudgal, C., & Vassileva, J. (2000). Bilateral negotiation with incomplete and uncertain information: a decision-theoretic approach using a model of the opponent. In *Proceedings of the 4th International Workshop on Cooperative Information Agents IV, The Future of Information Agents in Cyberspace* (pp. 107–118).
11. Noriega, P., & Sierra, C. (1998). Agent-mediated integrative negotiation for retail electronic commerce. In *Agent Mediated Electronic Commerce* (pp. 70–90). Minneapolis, Minnesota, May.
12. Parsons, S., Sierra, C., & Jennings, N. R. (1998). Agents that reason and negotiate by arguing. *Journal of Logic and Computation*, 8(3) 261–192.
13. Sandholm, T., & Vulkan, N. (1999). Bargaining with Deadlines. In *Proceedings of the National Conference on Artificial Intelligence (AAAI)* (pp. 44–51). Orlando, FL.
14. Teuteberg, F. (2003). Experimental evaluation of a model for multilateral negotiation with fuzzy preferences on an agent-based marketplace. *Electronic Markets*, 13(1), 21–32.
15. Zeng, D. & Sycara, K. (1998). Bayesian learning in negotiation. *International Journal of Human-Computer Studies*, 48, 125–141.